

SHAMIM MAHBUB

+88017-20506729 | sha237mim@gmail.com | Dhaka, Bangladesh

[GitHub](#) | [LinkedIn](#) | [Google Scholar](#)

Senior Generative AI & Agentic Systems Engineer with 5+ years building production-grade AI solutions. Expert in designing multi-agent orchestration systems, RAG architectures, and MCP-based tool integration. Proficient in LLM fine-tuning (LoRA/QLoRA), reasoning model evaluation (o3/DeepSeek-R1), multimodal AI, and LLMOps pipelines. Track record of delivering measurable business impact across enterprise clients in the US, Singapore, and Bangladesh.

PROFESSIONAL EXPERIENCE

GenAI Engineer — Artech Digital

Nov 2024 – Jan 2026

San Francisco, USA (Remote)

- Led end-to-end development of agentic AI systems for enterprise clients using LangGraph, CrewAI, and MCP-based tool integration, enabling autonomous multi-step task execution across complex workflows.
- Designed and deployed a WhatsApp-based e-commerce platform powered by LangGraph multi-agent orchestration; integrated M-Pesa Daraja APIs for fully automated end-to-end transactions.
- Architected a GEO Intelligence Platform for brand visibility tracking across generative search engines — built concurrent sampling pipelines, cost tracking, and response analytics for marketing teams.
- Built an Enterprise Bid Generation Agent (LangGraph + RAG) that auto-generates RFP responses by analyzing requirements and retrieving past proposals, reducing bid preparation time by 65%.
- Evaluated and benchmarked reasoning models (OpenAI o3, DeepSeek-R1) for client use cases, advising on model selection trade-offs between latency, cost, and accuracy.
- Mentored junior engineers on LLMOps best practices: model versioning, drift monitoring, prompt regression testing, and guardrail implementation against OWASP LLM Top 10 risks.

Machine Learning Engineer — Spekter GmbH

Jan 2024 – Nov 2024

Adabor, Dhaka, Bangladesh

- Developed anomaly detection models for industrial sensor data using time-series deep learning (LSTM, Transformer-based architectures), reducing false-positive alerts by ~30%.
- Implemented AWS CloudWatch monitoring pipelines and automated alerting for ML model performance degradation in production.

Software Engineer – AI & NLP — SkyRanko

Jan 2023 – Dec 2023

Rajshahi, Bangladesh

- Built NLP pipeline for paraphrasing and summarization of English text using fine-tuned T5/BART models.
- Developed automated web scraping systems for data collection pipelines feeding downstream ML models.

Software Engineer – NLP — Zibew E-Commerce PTE Ltd.

Mar 2022 – Dec 2022

Singapore (Remote)

- Developed and deployed AI chatbots for multiple client platforms; researched LLM-based conversational AI to replace rule-based systems.

Artificial Intelligence Engineer — Opus Technology Ltd.

Jul 2021 – Feb 2022

Dhaka, Bangladesh

- Built a Reinforcement Learning bot for CS:GO using PPO/DQN; developed company website chatbot using intent classification.

TECHNICAL SKILLS

Languages: Python (Expert), SQL, C (Intermediate)

Agentic & GenAI: LLM Fine-tuning, Prompt Engineering, RAG, Agentic RAG, MCP (Model Context Protocol), AI Agents, Multimodal AI, Reasoning Models (o3/o4, DeepSeek-R1, Gemini Thinking), LLMops, Guardrails / Red Teaming

Frameworks: LangChain, LangGraph, CrewAI, LlamaIndex, AutoGen (AG2), FastMCP, Keras, TensorFlow, PyTorch

Libraries: Transformers, FAISS, Sentence-Transformers, NLTK, SpaCy, NumPy, Pandas, OpenCV, Matplotlib, SciPy

Vector DBs: Pinecone, Chroma, Qdrant, Weaviate

Cloud & MLOps: AWS (SageMaker, Lambda, S3), Azure (OpenAI Service, ML Workspace), Vertex AI, MLflow, Weights & Biases, Docker, FastAPI, Flask, GitHub Actions

Tools & Platforms: HuggingFace, ComfyUI, Gradio, Streamlit, Selenium, Git & GitHub, MySQL, SQLServer, Linux

SELECTED PROJECTS

MCP-Integrated Enterprise RAG System — Built an MCP server (FastMCP) wrapping a domain-specific RAG pipeline, enabling portable tool integration across Claude, VS Code, and custom agent backends. Improved factual accuracy 47% vs. base LLM.

Reasoning Model Evaluation Framework — Designed a benchmarking system comparing o3, DeepSeek-R1, and Gemini Thinking models across enterprise use cases; delivered actionable model-selection guidance reducing inference costs by 30%.

Enterprise Bid Generation Agent — LangGraph multi-agent system that auto-generates customized RFP responses by analyzing requirement docs and retrieving past proposals; reduced bid preparation time by 65%.

GEO Intelligence Platform — Scalable LLM evaluation platform generating prompt variants, running multi-temperature sampling experiments, and analyzing AI brand visibility across generative search engines with cost tracking.

Medical Knowledge Assistant — Healthcare-focused RAG system using specialized medical embeddings and clinical knowledge bases; achieved 93% accuracy on physician-verified test cases.

LLM Fine-tuning Pipeline (LoRA/QLoRA) — Automated pipeline for parameter-efficient fine-tuning of open-source LLMs on proprietary data; reduced training costs by 68% vs. full fine-tuning.

WhatsApp Agentic Commerce Platform — LangGraph-powered conversational commerce on WhatsApp with AI-driven trade-in valuation, product recommendations, and integrated M-Pesa STK push payments — fully automated end-to-end.

PUBLICATIONS

- Hasan, M. J., **Mahbub, S.**, et al. "Rice disease identification and classification by integrating SVM with deep CNN." ICASERT 2019, IEEE. **(141 citations)**
- Wahid, M. F., Hasan, M. J., Alom, M. S., & **Mahbub, S.** "Performance Analysis of ML Techniques for Microscopic Bacteria Image Classification." ICCCNT 2019, IEEE. (15 citations)
- **Mahbub, S.**, et al. "Comparison of Handcrafted and Deep Neural Network Feature Extraction for Recognizing Bangla Sign Digits." ICCCNT 2021, IEEE.

EDUCATION

B.Sc. (Engg.) in Electrical & Electronic Engineering — Hajee Mohammad Danesh Science & Technology University (HSTU)

2016 – 2019

Dinajpur, Bangladesh

- CGPA: 3.34 / 4.00 | Research: Sleep Apnea detection via ECG (Deep Learning), Retinal Exudate Classification (DCNN)

RELEVANT CERTIFICATIONS & COURSEWORK

Advanced Generative AI Development • Building Production LLM Applications • RAG System Design & Implementation • LLM Fine-tuning & Model Alignment • LangChain & LlamaIndex for AI Applications • AI Agent Design & Orchestration • Vector Databases & Semantic Search • Multimodal AI Systems Development • Deep Learning Specialization • NLP Specialization • ChatGPT Prompt Engineering for Developers

LANGUAGES & LEADERSHIP

Languages: Bangla (Native), English (Professional) | **Leadership:** IEEE HSTU Student Branch – Membership Committee Convener; Joint Secretary, EEE Club of HSTU; Volunteer, National High School Programming Contest